

# CMDA 3634: mpiNonBlockingCircuitShell

May 29, 2018

## 1 Introduction

This document outlines the answers to the questions with written answers and Jumpshot images in homework 3 for CMDA 3634. This includes questions 3 and 4 in this assignment. Code for the homework is listed at the end of the document.

## 2 Q3

Check that your non-blocking MPI code reproduces the results of your serial code from HW01 to a reasonable tolerance. Use a convergence tolerance of  $1e-9$  to guarantee that the circuit currents have adequately converged for this test. Use the LATEX verbatim environment to report the loop currents output by your code.

From using the HW1 with a convergence tolerance of  $1e-9$ , I was able to get the following result.

```
epsilon = 1.00533e-09
epsilon = 1.00436e-09
epsilon = 1.00338e-09
epsilon = 1.00241e-09
epsilon = 1.00144e-09
epsilon = 1.00048e-09
epsilon = 9.99509e-10
I_{11} = 0.302347
I_{10 10} = 0.00317066
```

Using the same convergence tolerance, I was able to get almost identical result.

```
epsilon = 1.00482e-09
epsilon = 1.00385e-09
epsilon = 1.00288e-09
epsilon = 1.00191e-09
epsilon = 1.00094e-09
epsilon = 9.99973e-10
```

```
I_{11} = 0.302347
I_{10 10} = 0.00317066
```

### 3 Q4

Use jumpshot to visualize the time history of your code executing with four MPI processes on a circuit with total size 400x400. Describe any patterns you observe in the jumpshot traces. Remember to zoom in so that you can see individual MPI operations on the process timelines.

Figure 1, is the general overview of the jumpshot using the following executable command:

```
mpiexec -n 4 ./main 400 400 1e-6
```

As you look further in detail, figure 2, there are lots of MPI\_Allreduce and MPI\_Wait time spent among all 4 processors. Hard to tell whether MPI\_Isend or MPI\_Irecv are being any work done. Figure 3, this shows the detail of the time history work done by the process 2. As soon as the wait is finished, MPI\_Isend to which ever process that is designated to, then MPI\_Irecv happens.

### 4 Code

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <mpi.h>

/* function to convert from (i,j) cell index to linear storage index */
int idx(int N, int i, int j){

    return i + j*(N+2);
}

double allReduce(double data){
    // INPUT: data value from this MPI process to be summed up
    // OUTPUT: sum of the data variables from all MPI processes

    //int MPI_Allreduce(const void *sendbuf, void *recvbuf, int count,
    //                  MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)

    //int MPI_Iallreduce(const void *sendbuf, void *recvbuf, int count,
    //                  MPI_Datatype datatype, MPI_Op op, MPI_Comm comm, MPI_Request request)
    double someValue;
    //MPI_Comm comm = MPI_COMM_WORLD;
```

```

// int rank, size, i = 0;
//MPI.Comm_rank(MPLCOMM_WORLD, &rank);
// MPI.Comm_size(MPLCOMM_WORLD, &size);
// double *someData = (double*) calloc(size, sizeof(double));
//for (i = 0; i < )

//It is not strictly necessary to compute e at every iteration. Add a command li
//change the number of iterations between computation of e.

MPI_Allreduce(&data, &someValue, 1, MPLDOUBLE, MPLSUM, MPLCOMM_WORLD);

return someValue;

}

/* function to compute l2 norm (Euclidean norm of difference between Inew and Iold)
double calculateEpsilon(int M, int N, double *Iold, double *Inew){

double epsilon = 0;

int i, j;
for (j=1; j<=M; ++j){
    for (i=1; i<=N; ++i){
        epsilon += pow(Inew[idx(N, i, j)] - Iold[idx(N, i, j)], 2);
    }
}
epsilon = allReduce(epsilon);

// add up this quantity from all processes
epsilon = sqrt(epsilon);

return epsilon;
}

/* reference halo exchange from HW02 */
void haloExchange(int M, int N, double *I){

int rank, size, tag=999;
MPI_Status status;

MPI_Comm_rank(MPLCOMM_WORLD, &rank);
MPI_Comm_size(MPLCOMM_WORLD, &size);

if (rank < size - 1)
    MPI_Send(I + idx(N, 1, M), N, MPLDOUBLE, rank + 1, tag, MPLCOMM_WORLD);

```

```

    if (rank > 0)
        MPI_Send(I+idx(N,1,1), N, MPLDOUBLE, rank-1, tag, MPLCOMM_WORLD);

    if (rank > 0)
        MPI_Recv(I+idx(N,1,0), N, MPLDOUBLE, rank-1, tag, MPLCOMM_WORLD, &status);

    if (rank < size-1)
        MPI_Recv(I+idx(N,1,M+1), N, MPLDOUBLE, rank+1, tag, MPLCOMM_WORLD, &status);
}

```

```

void startHaloExchange(int M, int N, double *I, MPI_Request *IsendRequests, MPI_Request *IrecvRequests) {

    // initiate halo exchange using MPI_Isend and MPI_Irecv
    int rank, size, tag=999;
    MPI_Comm_rank(MPLCOMM_WORLD, &rank);
    MPI_Comm_size(MPLCOMM_WORLD, &size);
    MPI_Status status;

    // Your Q2a code starts here
    if (rank < size-1)
        MPI_Isend(I+idx(N,1,M), N, MPLDOUBLE, rank+1, tag, MPLCOMM_WORLD, IsendRequests+rank);
    if (rank > 0)
        MPI_Isend(I+idx(N,1,1), N, MPLDOUBLE, rank-1, tag, MPLCOMM_WORLD, IsendRequests+rank-1);
    if (rank > 0)
        MPI_Irecv(I+idx(N,1,0), N, MPLDOUBLE, rank-1, tag, MPLCOMM_WORLD, IrecvRequests+rank-1);
    if (rank < size-1)
        MPI_Irecv(I+idx(N,1,M+1), N, MPLDOUBLE, rank+1, tag, MPLCOMM_WORLD, IrecvRequests+rank+1);

    // Your Q2a code ends here
}

```

```

void endHaloRecv(MPI_Request *IrecvRequests) {

    // wait for halo data recv to complete using MPI_Wait
    int rank, size;
    MPI_Comm_rank(MPLCOMM_WORLD, &rank);
    MPI_Comm_size(MPLCOMM_WORLD, &size);

    // Your Q2b code starts here
}

```

```

    MPI_Status status;

    // MPI.Wait(IrecvRequests, &status);
    // Your Q2b code ends here

    if(rank < size - 1)
        MPI.Wait(IrecvRequests, &status);

    if(rank > 0)
        MPI.Wait(IrecvRequests + 1, &status);
}

void endHaloSend(MPI_Request *IsendRequests){

    // wait for outgoing halo data send to leave the buffer using MPI.Wait
    int rank, size;
    MPI_Comm_rank(MPLCOMM_WORLD, &rank);
    MPI_Comm_size(MPLCOMM_WORLD, &size);

    // Your Q2c code starts here
    MPI_Status status;

    //MPI.Wait(IsendRequests, &status);
    // Your Q2c code ends here

    if(rank < size - 1)
        MPI.Wait(IsendRequests, &status);

    if(rank > 0)
        MPI.Wait(IsendRequests + 1, &status);
}

/* function to update Inew from Iold */
void iterate(int M, int N, double *Iold, double *Inew){

    MPI_Request IsendRequests[2], IrecvRequests[2];
    int i, j;

    // int rank, size;
    // MPI_Comm_rank(MPLCOMM_WORLD, &rank);
    // MPI_Comm_size(MPLCOMM_WORLD, &size);
    // MPI_Status status;

    // initializes the swap of the top/bottom rows needed by the iterate step
    startHaloExchange(M, N, Iold, IsendRequests, IrecvRequests);

```

```

// process cell updates that do not require halo data from other processes
for (j=2;j<=M-1;++j)
    for (i=1;i<=N;++i)
        Inew[idx(N,i,j)] = 0.25*(Iold[idx(N,i+1,j)] + Iold[idx(N,i-1,j)] +
                                Iold[idx(N,i,j+1)] + Iold[idx(N,i,j-1)]);

// wait for the incoming halo data to arrive
endHaloRecv(IrecvRequests);

// finish update for bottom cells
j = 1;
for (i=1;i<=N;++i)
    Inew[idx(N,i,j)] = 0.25*(Iold[idx(N,i+1,j)] + Iold[idx(N,i-1,j)] +
                            Iold[idx(N,i,j+1)] + Iold[idx(N,i,j-1)]);

// finish update for top cells
j = M;
for (i=1;i<=N;++i)
    Inew[idx(N,i,j)] = 0.25*(Iold[idx(N,i+1,j)] + Iold[idx(N,i-1,j)] +
                            Iold[idx(N,i,j+1)] + Iold[idx(N,i,j-1)]);

// wait for the outgoing halo data buffer to be available for use
endHaloSend(IsendRequests);

}

```

```

/* function to solve for loop currents using Jacobi iterative method */
void solve(int M, int N, double tol){

    /* use for computed epsilon */
    double epsilon;

    double *Inew = (double*) calloc((M+2)*(N+2), sizeof(double));
    double *Iold = (double*) calloc((M+2)*(N+2), sizeof(double));

    /* set batteries based on MPI process rank*/
    int rank, size;
    MPI_Comm_rank(MPLCOMM_WORLD, &rank);
    MPI_Comm_size(MPLCOMM_WORLD, &size);
}

```

```

// Your Q1 code to set the ghost cells for the two batteries
// based on MPI process rank starts here
if(rank==0){ // bottom left cell
    Iold[idx(N,0,1)] = 1;
    Inew[idx(N,0,1)] = 1;
}

if(rank==size-1){ // top right cell
    printf(" cell\n");
    Iold[idx(N,N+1,M)] = 1;
    Inew[idx(N,N+1,M)] = 1;
}
// Your Q1 code to set the ghost cells ends here


/* iterate using the Jacobi method here */
do{

    /* iterate from Iold to Inew */
    iterate(M, N, Iold, Inew);

    /* iterate from Inew to Iold */
    iterate(M, N, Inew, Iold);

    /* compute epsilon (change in current) */
    epsilon = calculateEpsilon(M, N, Iold, Inew);

    /* print current residual */
    if(rank==0)
        printf("epsilon = %g\n", epsilon);

}while(epsilon>tol);

/* print out the loop current in cell (1 1) and (10 10) */
if(rank==0)
    printf("I_{11} = %g\n", Iold[idx(N,1,1)]);
if(rank*M <= 10 && 10 <= ((rank+1)*M))
    printf("I_{10 10} = %g\n", Iold[idx(N,10, 10-M*rank)]);
}

/* usage: ./main 100 1e-6
mpiexec -n 4 ./main 100 25 1e-6
to solve for a network of 100x100 to tolerance 1e-6 */

int main(int argc, char **argv){

```

```

// Your Q1 code to call MPI_Init starts here
MPI_Init(&argc, &argv);
// Your Q1 code to call MPI_Init ends here

{
    /* read N from the command line arguments */
    int N = atoi(argv[1]);

    /* read M from the command line arguments */
    int M = atoi(argv[2]);

    /* read the user supplied convergence from the command line arguments */
    double tol = atof(argv[3]);

    /* perform Jacobi iteration to solve for loop currents in resistor network */
    solve(M, N, tol);
}

// Your Q1 code to call MPI_Finalize starts here
MPI_Finalize();
// Your Q1 code to call MPI_Finalize ends here

return 0;
}

```



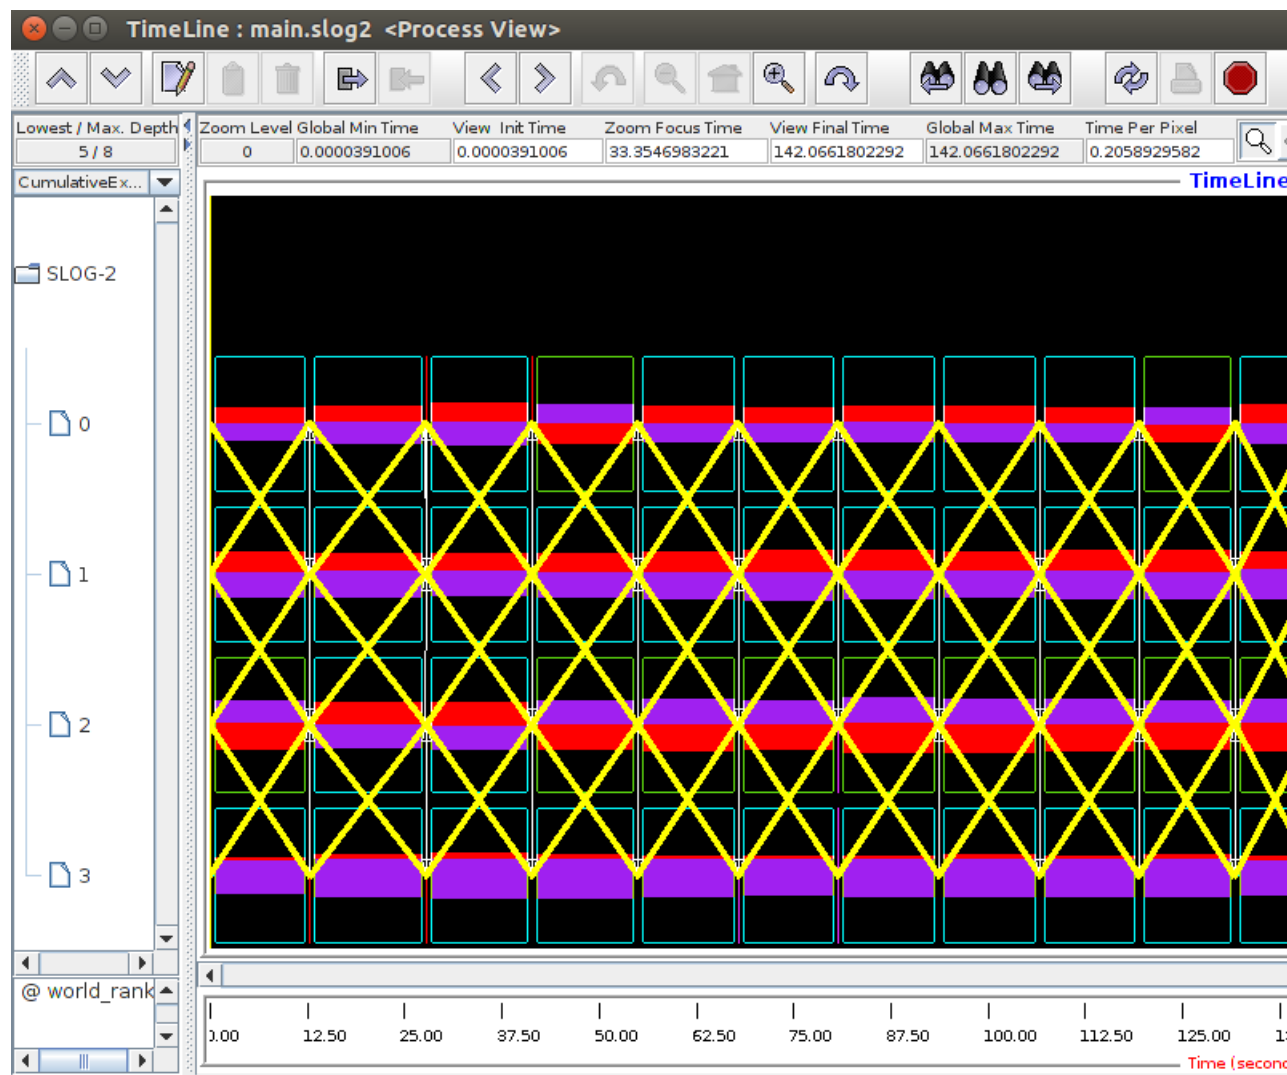


Figure 1: HW3 Jumpshot

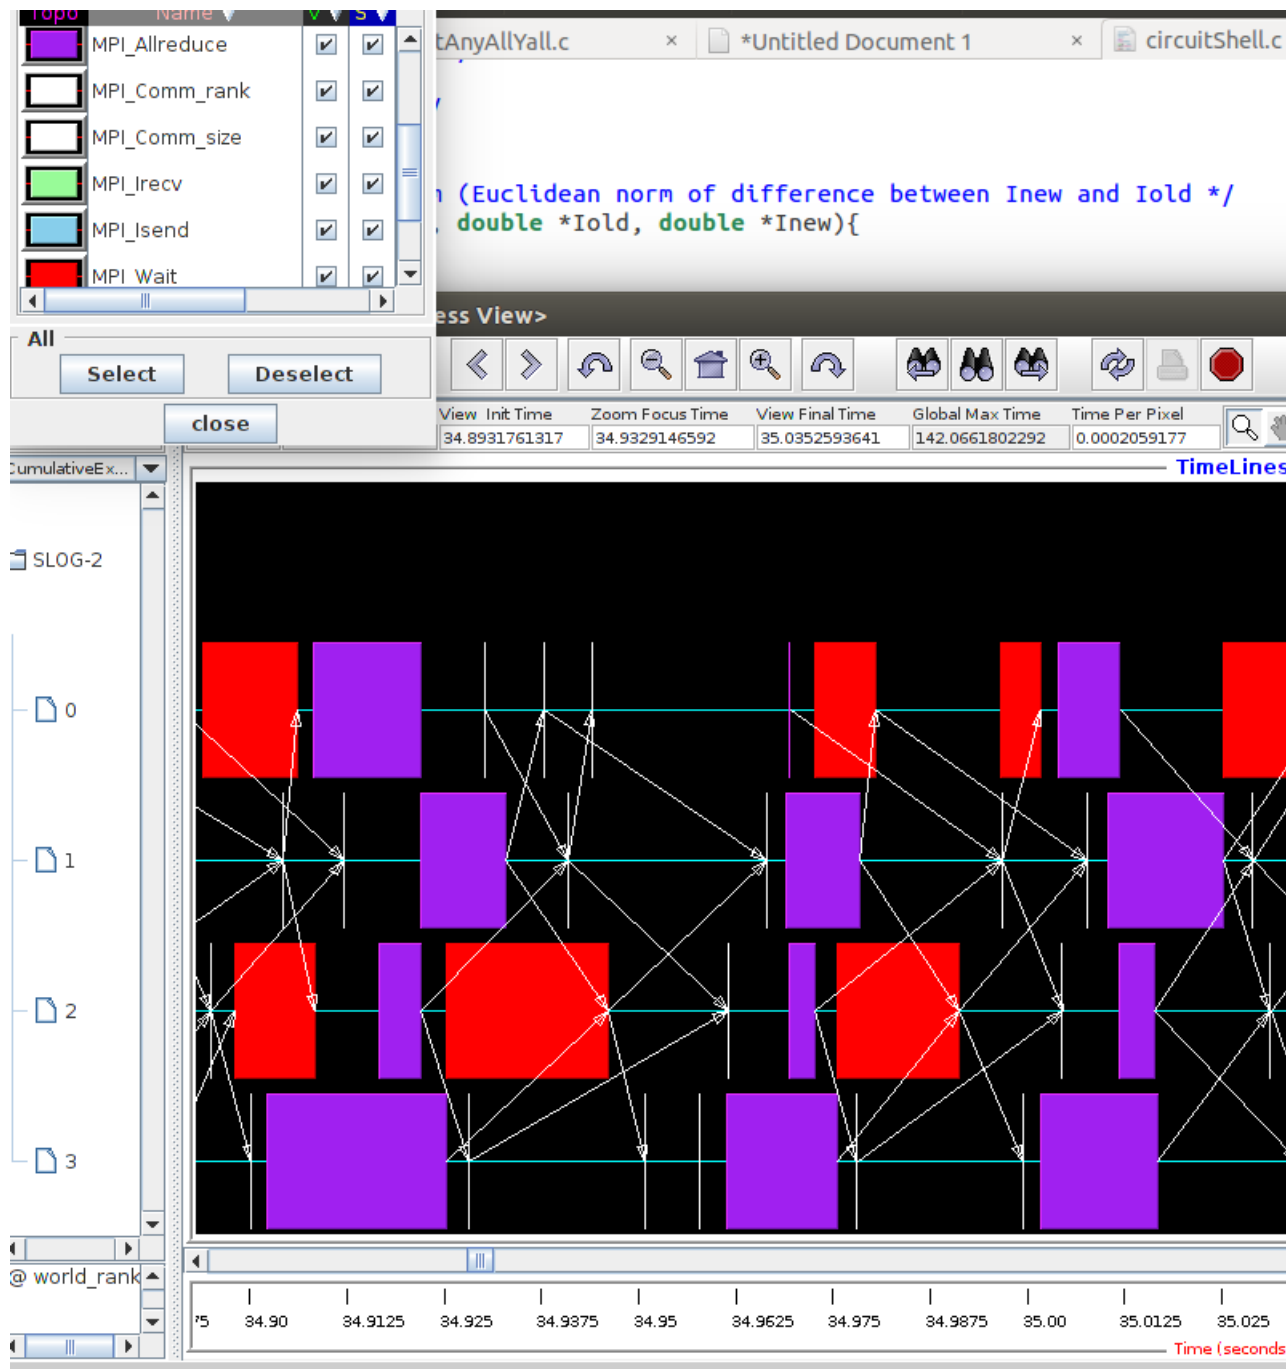


Figure 2: HW3 Jumpshot

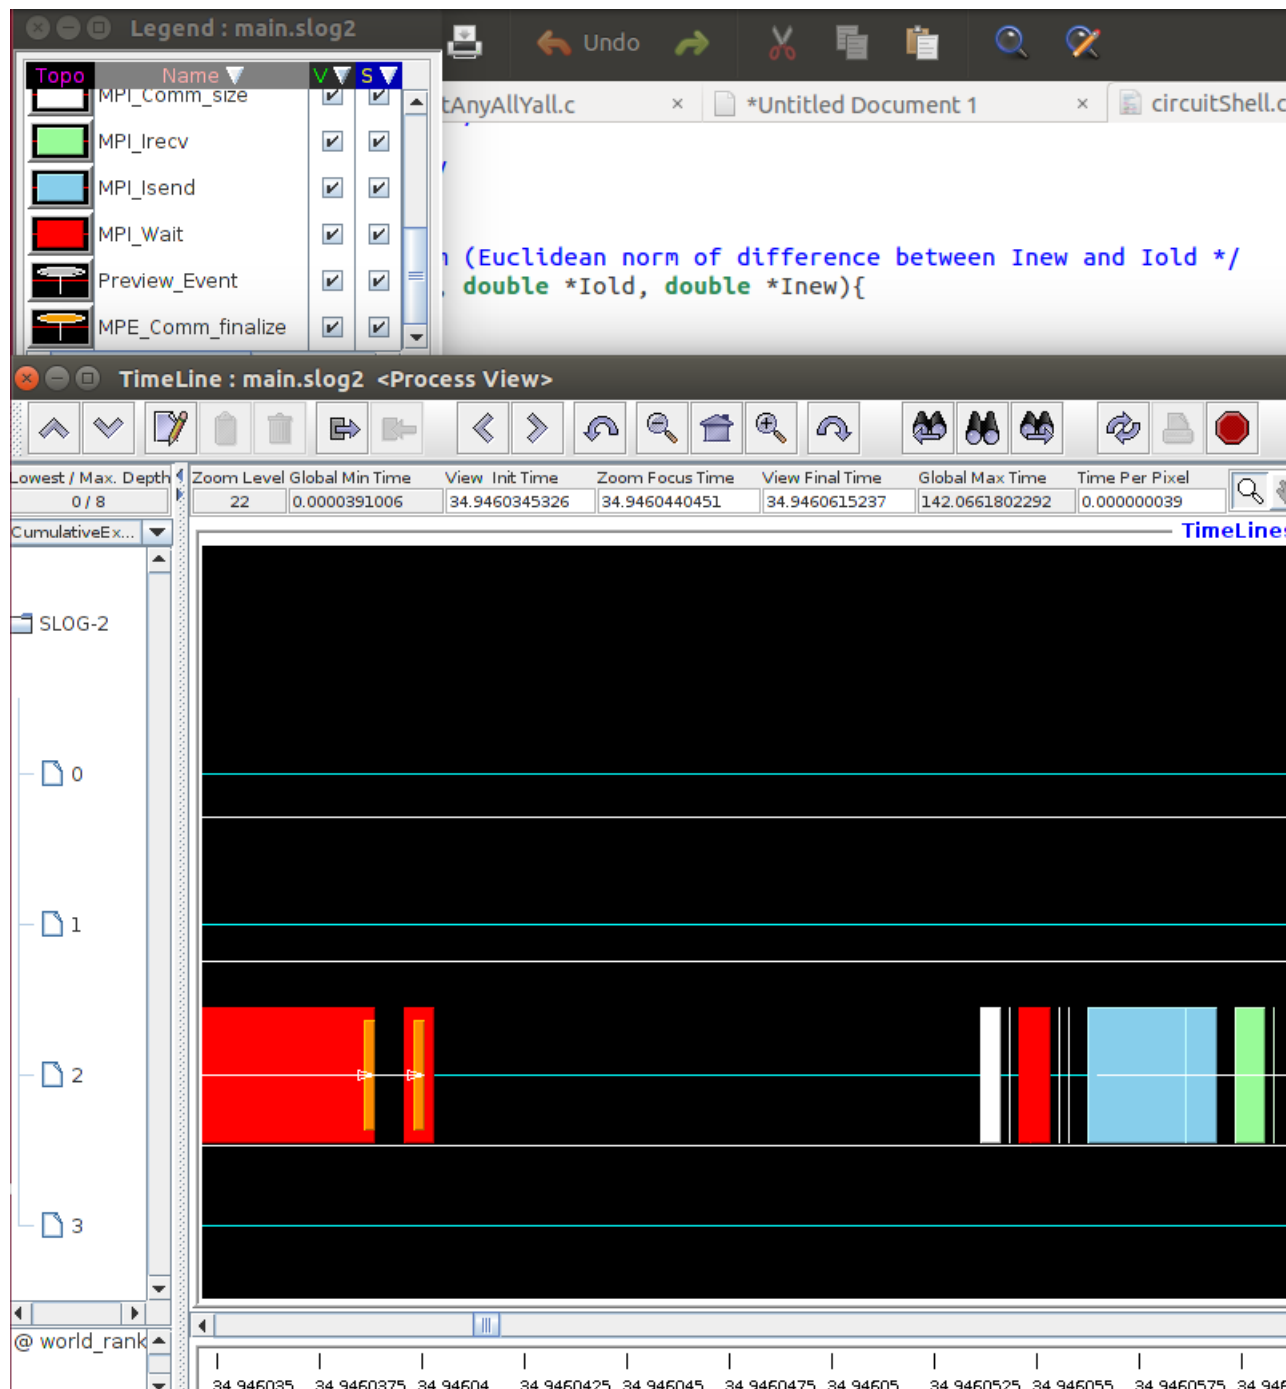


Figure 3: HW3 Jumpshot